

Convolutional Neural Networks and Computer Vision with TensorFlow

▼ Get the data

```
import zipfile

# Download zip file of pizza_steak images
!wget https://storage.googleapis.com/ztm\_tf\_course/food\_vision/pizza\_steak.zip

# Unzip the downloaded file
zip_ref = zipfile.ZipFile("pizza_steak.zip", "r")
zip_ref.extractall()
zip_ref.close()

--2023-05-11 03:50:16-- https://storage.googleapis.com/ztm\_tf\_course/food\_vision/pizza\_steak.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.128.128, 74.125.143.128, 173.194.69.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.128.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 109540975 (104M) [application/zip]
Saving to: 'pizza_steak.zip'

pizza_steak.zip      100%[=====] 104.47M  29.4MB/s    in 3.6s

2023-05-11 03:50:20 (29.4 MB/s) - 'pizza_steak.zip' saved [109540975/109540975]
```

▼ Inspect the data

```
!ls pizza_steak
```

```
test  train
```

Double-click (or enter) to edit

```
!ls pizza_steak/train/
```

```
pizza  steak
```

```
!ls pizza_steak/train/steak/
```

1000205.jpg	1647351.jpg	2238681.jpg	2824680.jpg	3375959.jpg	417368.jpg
100135.jpg	1650002.jpg	2238802.jpg	2825100.jpg	3381560.jpg	4176.jpg
101312.jpg	165639.jpg	2254705.jpg	2826987.jpg	3382936.jpg	42125.jpg
1021458.jpg	1658186.jpg	225990.jpg	2832499.jpg	3386119.jpg	421476.jpg
1032846.jpg	1658443.jpg	2260231.jpg	2832960.jpg	3388717.jpg	421561.jpg
10380.jpg	165964.jpg	2268692.jpg	285045.jpg	3389138.jpg	438871.jpg
1049459.jpg	167069.jpg	2271133.jpg	285147.jpg	3393547.jpg	43924.jpg
1053665.jpg	1675632.jpg	227576.jpg	2855315.jpg	3393688.jpg	440188.jpg
1068516.jpg	1678108.jpg	2283057.jpg	2856066.jpg	3396589.jpg	442757.jpg
1068975.jpg	168006.jpg	2286639.jpg	2859933.jpg	339891.jpg	443210.jpg
1081258.jpg	1682496.jpg	2287136.jpg	286219.jpg	3417789.jpg	444064.jpg
1090122.jpg	1684438.jpg	2291292.jpg	2862562.jpg	3425047.jpg	444709.jpg
1093966.jpg	168775.jpg	229323.jpg	2865730.jpg	3434983.jpg	447557.jpg
1098844.jpg	1697339.jpg	2300534.jpg	2878151.jpg	3435358.jpg	461187.jpg
1100074.jpg	1710569.jpg	2300845.jpg	2880035.jpg	3438319.jpg	461689.jpg
1105280.jpg	1714605.jpg	231296.jpg	2881783.jpg	3444407.jpg	465494.jpg
1117936.jpg	1724387.jpg	2315295.jpg	2884233.jpg	345734.jpg	468384.jpg
1126126.jpg	1724717.jpg	2323132.jpg	2890573.jpg	3460673.jpg	477486.jpg
114601.jpg	172936.jpg	2324994.jpg	2893832.jpg	3465327.jpg	482022.jpg
1147047.jpg	1736543.jpg	2327701.jpg	2893892.jpg	3466159.jpg	482465.jpg
1147883.jpg	1736968.jpg	2331076.jpg	2907177.jpg	3469024.jpg	483788.jpg
1155665.jpg	1746626.jpg	233964.jpg	290850.jpg	3470083.jpg	493029.jpg
1163977.jpg	1752330.jpg	2344227.jpg	2909831.jpg	3476564.jpg	503589.jpg
1190233.jpg	1761285.jpg	234626.jpg	2910418.jpg	3478318.jpg	510757.jpg
1208405.jpg	176508.jpg	234704.jpg	2912290.jpg	3488748.jpg	513129.jpg
1209120.jpg	1772039.jpg	2357281.jpg	2916448.jpg	3492328.jpg	513842.jpg
1212161.jpg	1777107.jpg	2361812.jpg	2916967.jpg	3518960.jpg	523535.jpg
1213988.jpg	1787505.jpg	2365287.jpg	2927833.jpg	3522209.jpg	525041.jpg
1219039.jpg	179293.jpg	2374582.jpg	2928643.jpg	3524429.jpg	534560.jpg
1225762.jpg	1816235.jpg	239025.jpg	2929179.jpg	3528458.jpg	534633.jpg
1230968.jpg	1822407.jpg	2390628.jpg	2936477.jpg	3531805.jpg	536535.jpg

```

1236155.jpg 1823263.jpg 2392910.jpg 2938012.jpg 3536023.jpg 541410.jpg
1241193.jpg 1826066.jpg 2394465.jpg 2938151.jpg 3538682.jpg 543691.jpg
1248337.jpg 1828502.jpg 2395127.jpg 2939678.jpg 3540750.jpg 560503.jpg
1257104.jpg 1828969.jpg 2396291.jpg 2940544.jpg 354329.jpg 561972.jpg
126345.jpg 1829045.jpg 2400975.jpg 2940621.jpg 3547166.jpg 56240.jpg
1264050.jpg 1829088.jpg 2403776.jpg 2949079.jpg 3553911.jpg 56409.jpg
1264154.jpg 1836332.jpg 2403907.jpg 295491.jpg 3556871.jpg 564530.jpg
1264858.jpg 1839025.jpg 240435.jpg 296268.jpg 355715.jpg 568972.jpg
127029.jpg 1839481.jpg 2404695.jpg 2964732.jpg 356234.jpg 576725.jpg
1289900.jpg 183995.jpg 2404884.jpg 2965021.jpg 3571963.jpg 588739.jpg
1290362.jpg 184110.jpg 2407770.jpg 2966859.jpg 3576078.jpg 590142.jpg
1295457.jpg 184226.jpg 2412263.jpg 2977966.jpg 3577618.jpg 60633.jpg
1312841.jpg 1846706.jpg 2425062.jpg 2979061.jpg 3577732.jpg 60655.jpg
1313316.jpg 1849364.jpg 2425389.jpg 2983260.jpg 3578934.jpg 606820.jpg
1324791.jpg 1849463.jpg 2435316.jpg 2984311.jpg 358042.jpg 612551.jpg
1327567.jpg 1849542.jpg 2437268.jpg 2988960.jpg 358045.jpg 614975.jpg
1327667.jpg 1853564.jpg 2437843.jpg 2989882.jpg 3591821.jpg 616809.jpg
1333055.jpg 1869467.jpg 2440131.jpg 2995169.jpg 359330.jpg 628628.jpg
1334054.jpg 1870942.jpg 2443168.jpg 2996324.jpg 3601483.jpg 632427.jpg
1335556.jpg 187303.jpg 2446660.jpg 3000131.jpg 3606642.jpg 636594.jpg
1337814.jpg 187521.jpg 2455944.jpg 3002350.jpg 3609394.jpg 637374.jpg
1340977.jpg 1888450.jpg 2458401.jpg 3007772.jpg 361067.jpg 640539.jpg
1343209.jpg 1889336.jpg 2487306.jpg 3008192.jpg 3613455.jpg 644777.jpg
134369.jpg 1907039.jpg 248841.jpg 3009617.jpg 3621464.jpg 644867.jpg
1344105.jpg 1925230.jpg 2489716.jpg 3011642.jpg 3621562.jpg 658189.jpg
134598.jpg 1927984.jpg 2490489.jpg 3020591.jpg 3621565.jpg 660900.jpg
1246207.jpg 1928577.jpg 2490591.jpg 3020579.jpg 3620555.jpg 662011.jpg

```

```

import os

for dirpath, dirnames, filenames in os.walk("pizza_steak"):
    print(f"There are {len(dirnames)} directories and {len(filenames)} images in '{dirpath}'.")

There are 2 directories and 0 images in 'pizza_steak'.
There are 2 directories and 0 images in 'pizza_steak/test'.
There are 0 directories and 250 images in 'pizza_steak/test/steak'.
There are 0 directories and 250 images in 'pizza_steak/test/pizza'.
There are 2 directories and 0 images in 'pizza_steak/train'.
There are 0 directories and 750 images in 'pizza_steak/train/steak'.
There are 0 directories and 750 images in 'pizza_steak/train/pizza'.

```

```

# Another way to find out how many images are in a file
num_steak_images_train = len(os.listdir("pizza_steak/train/steak"))

```

```
num_steak_images_train
```

```
750
```

```

# Get the class names
import pathlib
import numpy as np
data_dir = pathlib.Path("pizza_steak/train/")
class_names = np.array(sorted([item.name for item in data_dir.glob('*')])) # created a list of class_names from the subdirectories
print(class_names)

```

```
['pizza' 'steak']
```

```

# View an image
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import random

```

```

def view_random_image(target_dir, target_class):
    # Setup target directory (we'll view images from here)
    target_folder = target_dir+target_class

    # Get a random image path
    random_image = random.sample(os.listdir(target_folder), 1)

```

```

    # Read in the image and plot it using matplotlib
    img = mpimg.imread(target_folder + "/" + random_image[0])
    plt.imshow(img)
    plt.title(target_class)
    plt.axis("off");

```

```
print(f"Image shape: {img.shape}") # show the shape of the image
```

```
return img
```

```
# View a random image from the training dataset
img = view_random_image(target_dir="pizza_steak/train/",
                         target_class="steak")
```

Image shape: (512, 512, 3)

steak



```
# View the img (actually just a big array/tensor)
img
```

```
array([[[ 80,  32,  18],
       [ 77,  29,  15],
       [ 75,  27,  13],
       ...,
       [189, 163, 130],
       [196, 170, 137],
       [192, 166, 133]],

      [[ 79,  31,  17],
       [ 76,  28,  14],
       [ 74,  26,  12],
       ...,
       [155, 129,  96],
       [177, 151, 118],
       [192, 166, 133]],

      [[ 78,  30,  16],
       [ 76,  28,  14],
       [ 74,  26,  12],
       ...,
       [135, 109,  76],
       [149, 123,  90],
       [170, 144, 111]],

      ...,

      [[ 52,  18,  17],
       [ 63,  29,  28],
       [ 53,  19,  17],
       ...,
       [253, 237, 203],
       [253, 237, 203],
       [253, 237, 203]],

      [[ 53,  22,  20],
       [ 63,  32,  30],
       [ 56,  22,  21],
       ...,
       [253, 237, 203],
       [253, 237, 203],
       [252, 236, 202]],

      [[ 41,  10,   8],
       [ 62,  31,  29],
       [ 53,  19,  18],
       ...,
       [253, 237, 203],
       [252, 236, 202],
       [252, 236, 202]]], dtype=uint8)
```

```
# View the image shape
img.shape # returns (width, height, colour channels)

(512, 512, 3)

# Get all the pixel values between 0 & 1
img/255.

array([[[0.31372549, 0.1254902 , 0.07058824],
       [0.30196078, 0.11372549, 0.05882353],
       [0.29411765, 0.10588235, 0.05098039],
       ...,
       [0.74117647, 0.63921569, 0.50980392],
       [0.76862745, 0.66666667, 0.5372549 ],
       [0.75294118, 0.65098039, 0.52156863]],

      [[0.30980392, 0.12156863, 0.06666667],
       [0.29803922, 0.10980392, 0.05490196],
       [0.29019608, 0.10196078, 0.04705882],
       ...,
       [0.60784314, 0.50588235, 0.37647059],
       [0.69411765, 0.59215686, 0.4627451 ],
       [0.75294118, 0.65098039, 0.52156863]],

      [[0.30588235, 0.11764706, 0.0627451 ],
       [0.29803922, 0.10980392, 0.05490196],
       [0.29019608, 0.10196078, 0.04705882],
       ...,
       [0.52941176, 0.42745098, 0.29803922],
       [0.58431373, 0.48235294, 0.35294118],
       [0.66666667, 0.56470588, 0.43529412]],

      ...,

      [[0.20392157, 0.07058824, 0.06666667],
       [0.24705882, 0.11372549, 0.10980392],
       [0.20784314, 0.0745098 , 0.06666667],
       ...,
       [0.99215686, 0.92941176, 0.79607843],
       [0.99215686, 0.92941176, 0.79607843],
       [0.99215686, 0.92941176, 0.79607843]],

      [[0.20784314, 0.08627451, 0.07843137],
       [0.24705882, 0.1254902 , 0.11764706],
       [0.21960784, 0.08627451, 0.08235294],
       ...,
       [0.99215686, 0.92941176, 0.79607843],
       [0.99215686, 0.92941176, 0.79607843],
       [0.98823529, 0.9254902 , 0.79215686]],

      [[0.16078431, 0.03921569, 0.03137255],
       [0.24313725, 0.12156863, 0.11372549],
       [0.20784314, 0.0745098 , 0.07058824],
       ...,
       [0.99215686, 0.92941176, 0.79607843],
       [0.98823529, 0.9254902 , 0.79215686],
       [0.98823529, 0.9254902 , 0.79215686]]])

import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Set the seed
tf.random.set_seed(42)

# Preprocess data
train_datagen = ImageDataGenerator(rescale=1./255)
valid_datagen = ImageDataGenerator(rescale=1./255)

# Setup the train and test directories
train_dir = "pizza_steak/train/"
test_dir = "pizza_steak/test/"

# Import data from directories and turn it into batches
train_data = train_datagen.flow_from_directory(train_dir,
                                                batch_size=32, # number of images to process at a time
                                                target_size=(224, 224), # convert all images to be 224 x 224
                                                class_mode="binary", # type of problem
                                                seed=42)
```

```

valid_data = valid_datagen.flow_from_directory(test_dir,
                                              batch_size=32,
                                              target_size=(224, 224),
                                              class_mode="binary",
                                              seed=42)

model_1 = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(filters=10,
                          kernel_size=3, # can also be (3, 3)
                          activation="relu",
                          input_shape=(224, 224, 3)), # first layer specifies input shape (height, width, colour channels)
    tf.keras.layers.Conv2D(10, 3, activation="relu"),
    tf.keras.layers.MaxPool2D(pool_size=2, # pool_size can also be (2, 2)
                            padding="valid"), # padding can also be 'same'
    tf.keras.layers.Conv2D(10, 3, activation="relu"),
    tf.keras.layers.Conv2D(10, 3, activation="relu"),
    tf.keras.layers.MaxPool2D(2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(1, activation="sigmoid") # binary activation output
])

# Compile the model
model_1.compile(loss="binary_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=["accuracy"])

# Fit the model
history_1 = model_1.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=valid_data,
                        validation_steps=len(valid_data))

Found 1500 images belonging to 2 classes.
Found 500 images belonging to 2 classes.
Epoch 1/5
47/47 [=====] - 23s 209ms/step - loss: 0.5981 - accuracy: 0.6773 - val_loss: 0.4158 - val_accuracy: 0.8060
Epoch 2/5
47/47 [=====] - 9s 194ms/step - loss: 0.4584 - accuracy: 0.7967 - val_loss: 0.3675 - val_accuracy: 0.8500
Epoch 3/5
47/47 [=====] - 9s 192ms/step - loss: 0.4656 - accuracy: 0.7893 - val_loss: 0.4424 - val_accuracy: 0.7940
Epoch 4/5
47/47 [=====] - 9s 192ms/step - loss: 0.4098 - accuracy: 0.8380 - val_loss: 0.3485 - val_accuracy: 0.8780
Epoch 5/5
47/47 [=====] - 9s 191ms/step - loss: 0.3708 - accuracy: 0.8453 - val_loss: 0.3162 - val_accuracy: 0.8820

model_1.summary()

Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 222, 222, 10)	280
conv2d_1 (Conv2D)	(None, 220, 220, 10)	910
max_pooling2d (MaxPooling2D	(None, 110, 110, 10)	0
)		
conv2d_2 (Conv2D)	(None, 108, 108, 10)	910
conv2d_3 (Conv2D)	(None, 106, 106, 10)	910
max_pooling2d_1 (MaxPooling	(None, 53, 53, 10)	0
2D)		
flatten (Flatten)	(None, 28090)	0
dense (Dense)	(None, 1)	28091

```

=====
Total params: 31,101
Trainable params: 31,101
Non-trainable params: 0

```

```

# Set random seed
tf.random.set_seed(42)

```

```
# Create a model to replicate the TensorFlow model
model_2 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(224, 224, 3)), # dense layers expect a 1-dimensional vector as input
    tf.keras.layers.Dense(4, activation='relu'),
    tf.keras.layers.Dense(4, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
model_2.compile(loss='binary_crossentropy',
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=["accuracy"])

# Fit the model
history_2 = model_2.fit(train_data, # use same training data created above
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=valid_data, # use same validation data created above
                        validation_steps=len(valid_data))

Epoch 1/5
47/47 [=====] - 11s 197ms/step - loss: 1.7892 - accuracy: 0.4920 - val_loss: 0.6932 - val_accuracy: 0.5000
Epoch 2/5
47/47 [=====] - 9s 191ms/step - loss: 0.6932 - accuracy: 0.5000 - val_loss: 0.6932 - val_accuracy: 0.5000
Epoch 3/5
47/47 [=====] - 9s 189ms/step - loss: 0.6932 - accuracy: 0.5000 - val_loss: 0.6932 - val_accuracy: 0.5000
Epoch 4/5
47/47 [=====] - 9s 199ms/step - loss: 0.6932 - accuracy: 0.5000 - val_loss: 0.6931 - val_accuracy: 0.5000
Epoch 5/5
47/47 [=====] - 9s 196ms/step - loss: 0.6932 - accuracy: 0.5000 - val_loss: 0.6931 - val_accuracy: 0.5000

# second model's architecture
model_2.summary()

Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 150528)	0
dense_1 (Dense)	(None, 4)	602116
dense_2 (Dense)	(None, 4)	20
dense_3 (Dense)	(None, 1)	5

```
Total params: 602,141
Trainable params: 602,141
Non-trainable params: 0
```

```
# Set random seed
tf.random.set_seed(42)

# Create a model similar to model_1 but add an extra layer and increase the number of hidden units in each layer
model_3 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(224, 224, 3)), # dense layers expect a 1-dimensional vector as input
    tf.keras.layers.Dense(100, activation='relu'), # increase number of neurons from 4 to 100 (for each layer)
    tf.keras.layers.Dense(100, activation='relu'),
    tf.keras.layers.Dense(100, activation='relu'), # add an extra layer
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
model_3.compile(loss='binary_crossentropy',
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=["accuracy"])

# Fit the model
history_3 = model_3.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=valid_data,
                        validation_steps=len(valid_data))

Epoch 1/5
47/47 [=====] - 11s 201ms/step - loss: 2.8247 - accuracy: 0.6413 - val_loss: 0.6214 - val_accuracy: 0.7720
```

```

Epoch 2/5
47/47 [=====] - 9s 195ms/step - loss: 0.6881 - accuracy: 0.7300 - val_loss: 0.5762 - val_accuracy: 0.7140
Epoch 3/5
47/47 [=====] - 9s 196ms/step - loss: 1.2195 - accuracy: 0.6860 - val_loss: 0.6234 - val_accuracy: 0.7320
Epoch 4/5
47/47 [=====] - 9s 194ms/step - loss: 0.5172 - accuracy: 0.7893 - val_loss: 0.5061 - val_accuracy: 0.7600
Epoch 5/5
47/47 [=====] - 9s 192ms/step - loss: 0.5041 - accuracy: 0.7833 - val_loss: 0.4380 - val_accuracy: 0.8000

```

Woah! Looks like our model is learning again. It got ~70% accuracy on the training set and ~70% accuracy on the validation set.

How does the architecture look?

```
# Check out model_3 architecture
model_3.summary()
```

```

Model: "sequential_2"
-----
```

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 150528)	0
dense_4 (Dense)	(None, 100)	15052900
dense_5 (Dense)	(None, 100)	10100
dense_6 (Dense)	(None, 100)	10100
dense_7 (Dense)	(None, 1)	101

```

-----
```

```
Total params: 15,073,201
Trainable params: 15,073,201
Non-trainable params: 0
```

```
# import zipfile
# # Download zip file of pizza_steak images
# !wget https://storage.googleapis.com/ztm_tf_course/food_vision/pizza_steak.zip
```

```
# # Unzip the downloaded file
# zip_ref = zipfile.ZipFile("pizza_steak.zip", "r")
# zip_ref.extractall()
# zip_ref.close()
```

```
plt.figure()
plt.subplot(1, 2, 1)
steak_img = view_random_image("pizza_steak/train/", "steak")
plt.subplot(1, 2, 2)
pizza_img = view_random_image("pizza_steak/train/", "pizza")
```

```
Image shape: (512, 512, 3)
Image shape: (512, 512, 3)
```



```
# Define training and test directory paths
train_dir = "pizza_steak/train/"
test_dir = "pizza_steak/test/"
```

```
# Create train and test data generators and rescale the data
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(rescale=1/255.)
test_datagen = ImageDataGenerator(rescale=1/255.)

# Turn it into batches
train_data = train_datagen.flow_from_directory(directory=train_dir,
                                                target_size=(224, 224),
                                                class_mode='binary',
                                                batch_size=32)

test_data = test_datagen.flow_from_directory(directory=test_dir,
                                              target_size=(224, 224),
                                              class_mode='binary',
                                              batch_size=32)

Found 1500 images belonging to 2 classes.
Found 500 images belonging to 2 classes.

# Get a sample of the training data batch
images, labels = train_data.next() # get the 'next' batch of images/labels
len(images), len(labels)

(32, 32)

# Get the first two images
images[:2], images[0].shape

(array([[[[0.47058827, 0.40784317, 0.34509805],
          [0.4784314 , 0.427451 , 0.3647059 ],
          [0.48627454, 0.43529415, 0.37254903],
          ...,
          [0.8313726 , 0.70980394, 0.48627454],
          [0.8431373 , 0.73333335, 0.5372549 ],
          [0.87843144, 0.7725491 , 0.5882353 ]],  
[[[0.50980395, 0.427451 , 0.36078432],
          [0.5058824 , 0.42352945, 0.35686275],
          [0.5137255 , 0.4431373 , 0.3647059 ],
          ...,
          [0.82745105, 0.7058824 , 0.48235297],
          [0.82745105, 0.70980394, 0.5058824 ],
          [0.8431373 , 0.73333335, 0.5372549 ]],  
[[[0.5254902 , 0.427451 , 0.34901962],
          [0.5372549 , 0.43921572, 0.36078432],
          [0.5372549 , 0.45098042, 0.36078432],
          ...,
          [0.82745105, 0.7019608 , 0.4784314 ],
          [0.82745105, 0.7058824 , 0.49411768],
          [0.8352942 , 0.7176471 , 0.5137255 ]],  
...,
[[[0.77647066, 0.5647059 , 0.2901961 ],
          [0.7803922 , 0.53333336, 0.22352943],
          [0.79215693, 0.5176471 , 0.18039216],
          ...,
          [0.30588236, 0.2784314 , 0.24705884],
          [0.24705884, 0.23137257, 0.19607845],
          [0.2784314 , 0.27450982, 0.25490198]]],  
[[[0.7843138 , 0.57254905, 0.29803923],
          [0.79215693, 0.54509807, 0.24313727],
          [0.8000001 , 0.5254902 , 0.18823531],
          ...,
          [0.2627451 , 0.23529413, 0.20392159],
          [0.24313727, 0.227451 , 0.19215688],
          [0.26666668, 0.2627451 , 0.24313727]]],  
[[[0.7960785 , 0.59607846, 0.3372549 ],
          [0.7960785 , 0.5647059 , 0.26666668],
          [0.81568635, 0.54901963, 0.22352943],
          ...,
          [0.23529413, 0.19607845, 0.16078432],
          [0.3019608 , 0.26666668, 0.24705884],
          [0.26666668, 0.2509804 , 0.24705884]]],  
[[[0.38823533, 0.4666667 , 0.36078432],
          [0.3921569 , 0.46274513, 0.36078432],
          [0.38431376, 0.454902 , 0.36078432],
```

```

...,
[0.5294118 , 0.627451 , 0.54509807],
[0.5294118 , 0.627451 , 0.54509807],
[0.5411765 . 0.6392157 . 0.5568628 ]]

# View the first batch of labels
labels

array([1., 1., 0., 0., 0., 1., 0., 1., 0., 0., 1., 0., 0., 0., 1.,
       1., 0., 1., 0., 1., 1., 0., 0., 0., 0., 1., 0., 1.],
      dtype=float32)

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPool2D, Activation
from tensorflow.keras import Sequential

# Create the model
model_4 = Sequential([
    Conv2D(filters=10,
           kernel_size=3,
           strides=1,
           padding='valid',
           activation='relu',
           input_shape=(224, 224, 3)), # input layer
    Conv2D(10, 3, activation='relu'),
    Conv2D(10, 3, activation='relu'),
    Flatten(),
    Dense(1, activation='sigmoid') # output layer
])

model_4.compile(loss='binary_crossentropy',
                 optimizer=Adam(),
                 metrics=['accuracy'])

len(train_data), len(test_data)

(47, 16)

# Fit the model
history_4 = model_4.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=test_data,
                        validation_steps=len(test_data))

Epoch 1/5
47/47 [=====] - 11s 198ms/step - loss: 0.6262 - accuracy: 0.6733 - val_loss: 0.4172 - val_accuracy: 0.8140
Epoch 2/5
47/47 [=====] - 9s 192ms/step - loss: 0.4260 - accuracy: 0.8253 - val_loss: 0.3766 - val_accuracy: 0.8420
Epoch 3/5
47/47 [=====] - 9s 191ms/step - loss: 0.2717 - accuracy: 0.9100 - val_loss: 0.3451 - val_accuracy: 0.8580
Epoch 4/5
47/47 [=====] - 9s 191ms/step - loss: 0.1141 - accuracy: 0.9720 - val_loss: 0.3705 - val_accuracy: 0.8440
Epoch 5/5
47/47 [=====] - 9s 191ms/step - loss: 0.0478 - accuracy: 0.9900 - val_loss: 0.5217 - val_accuracy: 0.8220

```

▼ 5. Evaluate the model

```

# Plot the training curves
import pandas as pd
pd.DataFrame(history_4.history).plot(figsize=(10, 7));

```



```
# Plot the validation and training data separately
def plot_loss_curves(history):
    """
    Returns separate loss curves for training and validation metrics.
    """
    loss = history.history['loss']
    val_loss = history.history['val_loss']

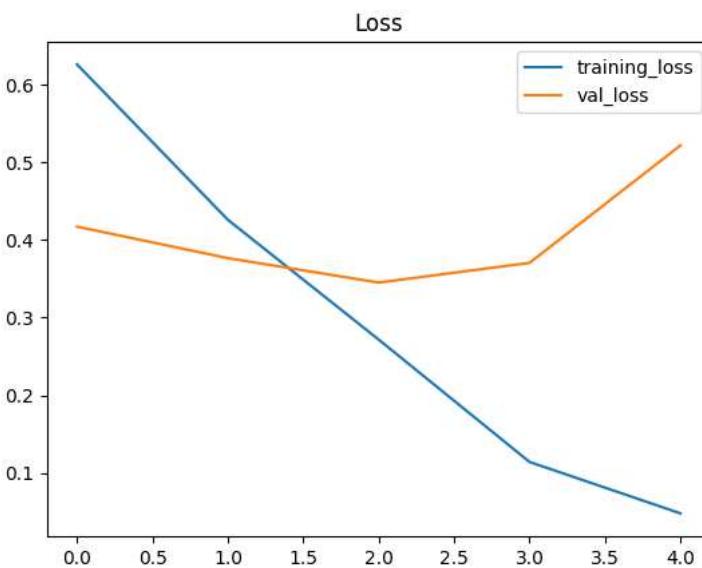
    accuracy = history.history['accuracy']
    val_accuracy = history.history['val_accuracy']

    epochs = range(len(history.history['loss']))

    # Plot loss
    plt.plot(epochs, loss, label='training_loss')
    plt.plot(epochs, val_loss, label='val_loss')
    plt.title('Loss')
    plt.xlabel('Epochs')
    plt.legend()

    # Plot accuracy
    plt.figure()
    plt.plot(epochs, accuracy, label='training_accuracy')
    plt.plot(epochs, val_accuracy, label='val_accuracy')
    plt.title('Accuracy')
    plt.xlabel('Epochs')
    plt.legend();

plot_loss_curves(history_4)
```



```
# model's architecture
model_4.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_4 (Conv2D)	(None, 222, 222, 10)	280
conv2d_5 (Conv2D)	(None, 220, 220, 10)	910
conv2d_6 (Conv2D)	(None, 218, 218, 10)	910
flatten_3 (Flatten)	(None, 475240)	0
dense_8 (Dense)	(None, 1)	475241
<hr/>		
Total params:	477,341	
Trainable params:	477,341	
Non-trainable params:	0	

```
# Create the model
model_5 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)),
    MaxPool2D(pool_size=2), # reduce number of features by half
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(1, activation='sigmoid')
])
```

```
# Compile model
model_5.compile(loss='binary_crossentropy',
                  optimizer=Adam(),
                  metrics=['accuracy'])
```

```
# Fit the model
history_5 = model_5.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=test_data,
                        validation_steps=len(test_data))
```

```
Epoch 1/5
47/47 [=====] - 11s 196ms/step - loss: 0.6482 - accuracy: 0.6240 - val_loss: 0.5451 - val_accuracy: 0.7100
Epoch 2/5
47/47 [=====] - 9s 190ms/step - loss: 0.4537 - accuracy: 0.8013 - val_loss: 0.3536 - val_accuracy: 0.8380
Epoch 3/5
47/47 [=====] - 9s 192ms/step - loss: 0.4161 - accuracy: 0.8233 - val_loss: 0.3578 - val_accuracy: 0.8460
Epoch 4/5
```

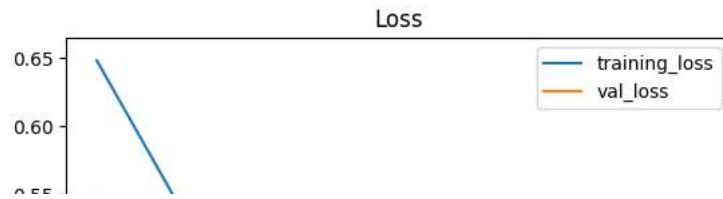
```
47/47 [=====] - 9s 192ms/step - loss: 0.3903 - accuracy: 0.8353 - val_loss: 0.3509 - val_accuracy: 0.8480
Epoch 5/5
47/47 [=====] - 9s 192ms/step - loss: 0.3567 - accuracy: 0.8513 - val_loss: 0.3200 - val_accuracy: 0.8700
```

```
model_5.summary()
```

```
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_7 (Conv2D)	(None, 222, 222, 10)	280
max_pooling2d_2 (MaxPooling 2D)	(None, 111, 111, 10)	0
conv2d_8 (Conv2D)	(None, 109, 109, 10)	910
max_pooling2d_3 (MaxPooling 2D)	(None, 54, 54, 10)	0
conv2d_9 (Conv2D)	(None, 52, 52, 10)	910
max_pooling2d_4 (MaxPooling 2D)	(None, 26, 26, 10)	0
flatten_4 (Flatten)	(None, 6760)	0
dense_9 (Dense)	(None, 1)	6761
<hr/>		
Total params: 8,861		
Trainable params: 8,861		
Non-trainable params: 0		

```
# Plot loss curves
plot_loss_curves(history_5)
```



```
# Create ImageDataGenerator training instance with data augmentation
train_datagen_augmented = ImageDataGenerator(rescale=1/255.,
                                             rotation_range=20, # rotate the image slightly between 0 and 20 degrees (note: this is an int not a float)
                                             shear_range=0.2, # shear the image
                                             zoom_range=0.2, # zoom into the image
                                             width_shift_range=0.2, # shift the image width ways
                                             height_shift_range=0.2, # shift the image height ways
                                             horizontal_flip=True) # flip the image on the horizontal axis

# Create ImageDataGenerator training instance without data augmentation
train_datagen = ImageDataGenerator(rescale=1/255.)

# Create ImageDataGenerator test instance without data augmentation
test_datagen = ImageDataGenerator(rescale=1/255.)
    epochs

# Import data and augment it from training directory
print("Augmented training images:")
train_data_augmented = train_datagen_augmented.flow_from_directory(train_dir,
                                                                    target_size=(224, 224),
                                                                    batch_size=32,
                                                                    class_mode='binary',
                                                                    shuffle=False)

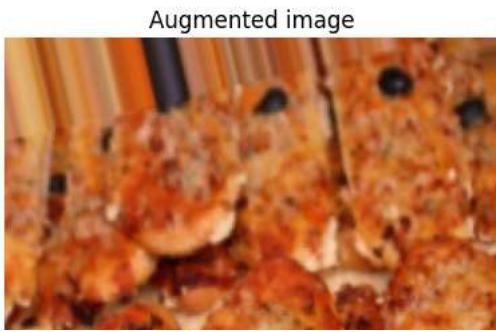
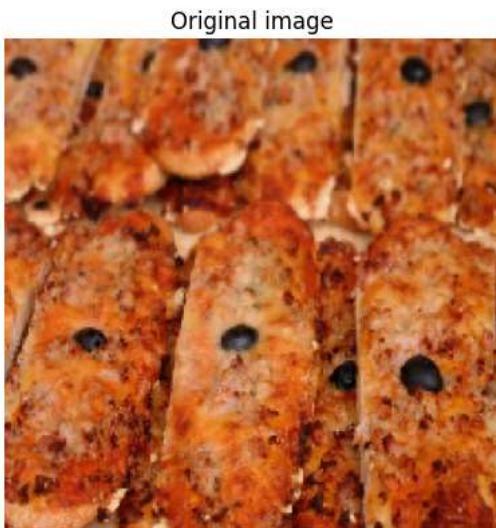
# Create non-augmented data batches
print("Non-augmented training images:")
train_data = train_datagen.flow_from_directory(train_dir,
                                               target_size=(224, 224),
                                               batch_size=32,
                                               class_mode='binary',
                                               shuffle=False)

print("Unchanged test images:")
test_data = test_datagen.flow_from_directory(test_dir,
                                              target_size=(224, 224),
                                              batch_size=32,
                                              class_mode='binary')

Augmented training images:
Found 1500 images belonging to 2 classes.
Non-augmented training images:
Found 1500 images belonging to 2 classes.
Unchanged test images:
Found 500 images belonging to 2 classes.

# Get data batch samples
images, labels = train_data.next()
augmented_images, augmented_labels = train_data_augmented.next()

# Show original image and augmented image
random_number = random.randint(0, 31)
plt.imshow(images[random_number])
plt.title(f"Original image")
plt.axis(False)
plt.figure()
plt.imshow(augmented_images[random_number])
plt.title(f"Augmented image")
plt.axis(False);
```



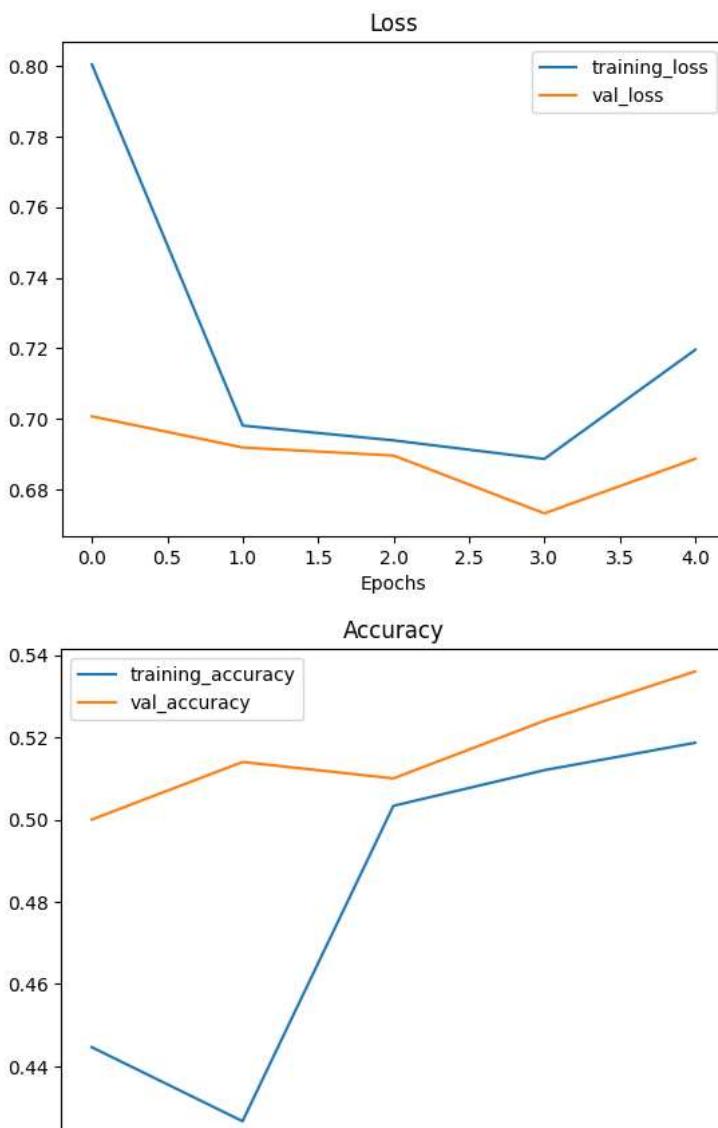
```
# Create the model
model_6 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)),
    MaxPool2D(pool_size=2), # reduce number of features by half
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(1, activation='sigmoid')
])

# Compile the model
model_6.compile(loss='binary_crossentropy',
                 optimizer=Adam(),
                 metrics=['accuracy'])

# Fit the model
history_6 = model_6.fit(train_data_augmented,
                        epochs=5,
                        steps_per_epoch=len(train_data_augmented),
                        validation_data=test_data,
                        validation_steps=len(test_data))

Epoch 1/5
47/47 [=====] - 24s 472ms/step - loss: 0.8005 - accuracy: 0.4447 - val_loss: 0.7007 - val_accuracy: 0.5000
Epoch 2/5
47/47 [=====] - 22s 475ms/step - loss: 0.6981 - accuracy: 0.4267 - val_loss: 0.6918 - val_accuracy: 0.5140
Epoch 3/5
47/47 [=====] - 22s 474ms/step - loss: 0.6938 - accuracy: 0.5033 - val_loss: 0.6896 - val_accuracy: 0.5100
Epoch 4/5
47/47 [=====] - 22s 477ms/step - loss: 0.6886 - accuracy: 0.5120 - val_loss: 0.6732 - val_accuracy: 0.5240
Epoch 5/5
47/47 [=====] - 22s 477ms/step - loss: 0.7195 - accuracy: 0.5187 - val_loss: 0.6886 - val_accuracy: 0.5360

# Check model's performance history training on augmented data
plot_loss_curves(history_6)
```



```
# Import data and augment it from directories
train_data_augmented_shuffled = train_datagen_augmented.flow_from_directory(train_dir,
                                                               target_size=(224, 224),
                                                               batch_size=32,
                                                               class_mode='binary',
                                                               shuffle=True) # Shuffle data (default)
```

Found 1500 images belonging to 2 classes.

```
# Create the model (same as model_5 and model_6)
model_7 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(1, activation='sigmoid')
])

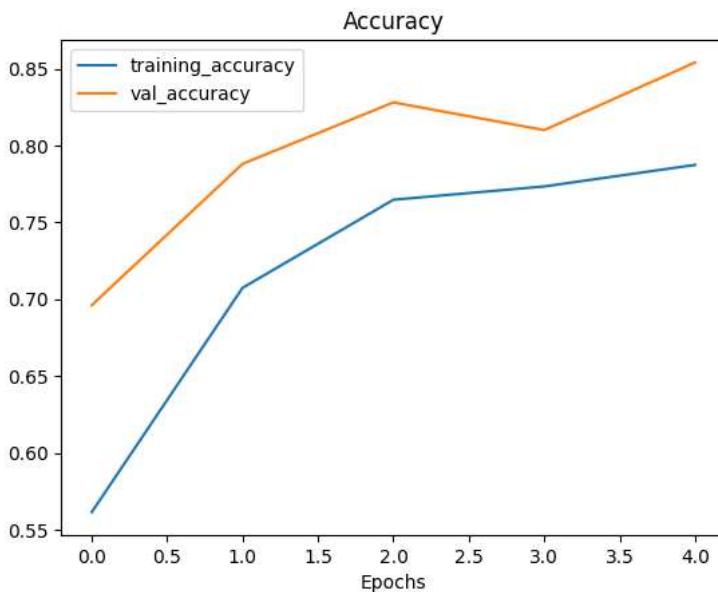
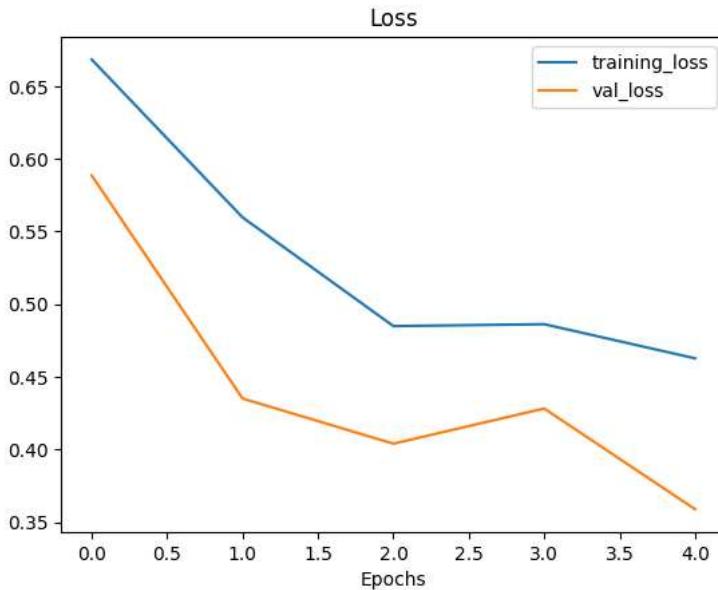
# Compile the model
model_7.compile(loss='binary_crossentropy',
                 optimizer=Adam(),
                 metrics=['accuracy'])

# Fit the model
history_7 = model_7.fit(train_data_augmented_shuffled, # now the augmented data is shuffled
                        epochs=5,
                        steps_per_epoch=len(train_data_augmented_shuffled),
```

```
validation_data=test_data,
validation_steps=len(test_data))
```

```
Epoch 1/5
47/47 [=====] - 24s 474ms/step - loss: 0.6685 - accuracy: 0.5613 - val_loss: 0.5886 - val_accuracy: 0.6960
Epoch 2/5
47/47 [=====] - 22s 478ms/step - loss: 0.5598 - accuracy: 0.7073 - val_loss: 0.4351 - val_accuracy: 0.7880
Epoch 3/5
47/47 [=====] - 22s 477ms/step - loss: 0.4849 - accuracy: 0.7647 - val_loss: 0.4040 - val_accuracy: 0.8280
Epoch 4/5
47/47 [=====] - 23s 487ms/step - loss: 0.4862 - accuracy: 0.7733 - val_loss: 0.4282 - val_accuracy: 0.8100
Epoch 5/5
47/47 [=====] - 23s 486ms/step - loss: 0.4627 - accuracy: 0.7873 - val_loss: 0.3589 - val_accuracy: 0.8540
```

```
# Check model's performance history training on augmented data
plot_loss_curves(history_7)
```



```
# Create a CNN model
model_8 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(1, activation='sigmoid')
])
```

```
# Compile the model
model_8.compile(loss="binary_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=["accuracy"])

# Fit the model
history_8 = model_8.fit(train_data_augmented_shuffled,
                        epochs=5,
                        steps_per_epoch=len(train_data_augmented_shuffled),
                        validation_data=test_data,
                        validation_steps=len(test_data))

Epoch 1/5
47/47 [=====] - 25s 492ms/step - loss: 0.6345 - accuracy: 0.6267 - val_loss: 0.4843 - val_accuracy: 0.7500
Epoch 2/5
47/47 [=====] - 23s 495ms/step - loss: 0.5293 - accuracy: 0.7327 - val_loss: 0.3829 - val_accuracy: 0.8400
Epoch 3/5
47/47 [=====] - 23s 490ms/step - loss: 0.5296 - accuracy: 0.7360 - val_loss: 0.5177 - val_accuracy: 0.7160
Epoch 4/5
47/47 [=====] - 22s 475ms/step - loss: 0.4971 - accuracy: 0.7560 - val_loss: 0.3477 - val_accuracy: 0.8600
Epoch 5/5
47/47 [=====] - 22s 476ms/step - loss: 0.4670 - accuracy: 0.7927 - val_loss: 0.3511 - val_accuracy: 0.8540
```

```
# Check model_1 architecture
model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 222, 222, 10)	280
conv2d_1 (Conv2D)	(None, 220, 220, 10)	910
max_pooling2d (MaxPooling2D	(None, 110, 110, 10)	0
)		
conv2d_2 (Conv2D)	(None, 108, 108, 10)	910
conv2d_3 (Conv2D)	(None, 106, 106, 10)	910
max_pooling2d_1 (MaxPooling	(None, 53, 53, 10)	0
2D)		
flatten (Flatten)	(None, 28090)	0
dense (Dense)	(None, 1)	28091
<hr/>		
Total params:	31,101	
Trainable params:	31,101	
Non-trainable params:	0	

```
# Check model_8 architecture
model_8.summary()
```

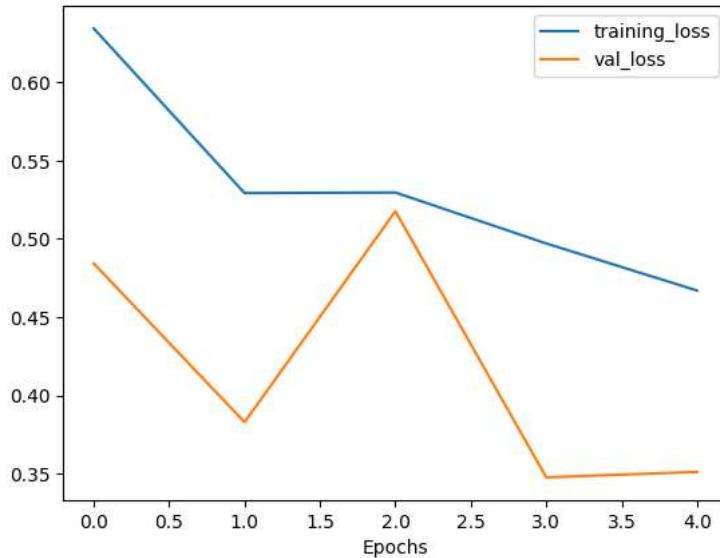
Model: "sequential_7"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_16 (Conv2D)	(None, 222, 222, 10)	280
conv2d_17 (Conv2D)	(None, 220, 220, 10)	910
max_pooling2d_11 (MaxPooling2D)	(None, 110, 110, 10)	0
conv2d_18 (Conv2D)	(None, 108, 108, 10)	910
conv2d_19 (Conv2D)	(None, 106, 106, 10)	910
max_pooling2d_12 (MaxPooling2D)	(None, 53, 53, 10)	0
flatten_7 (Flatten)	(None, 28090)	0
dense_12 (Dense)	(None, 1)	28091
<hr/>		

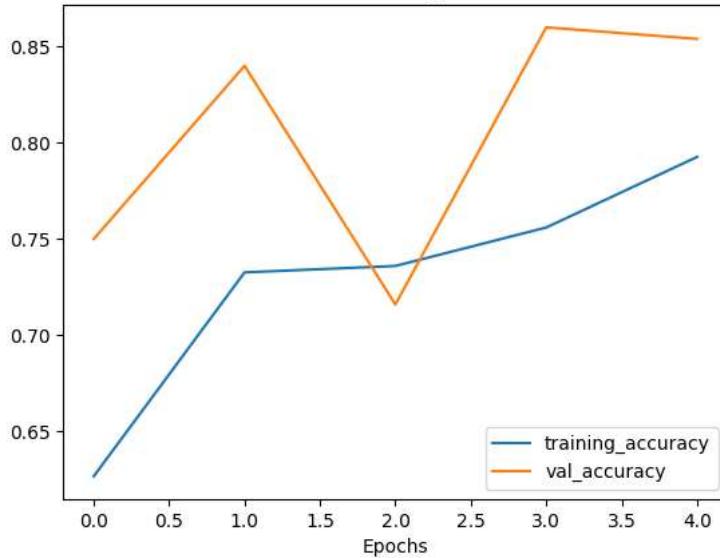
```
Total params: 31,101  
Trainable params: 31,101  
Non-trainable params: 0
```

```
plot_loss_curves(history_8)
```

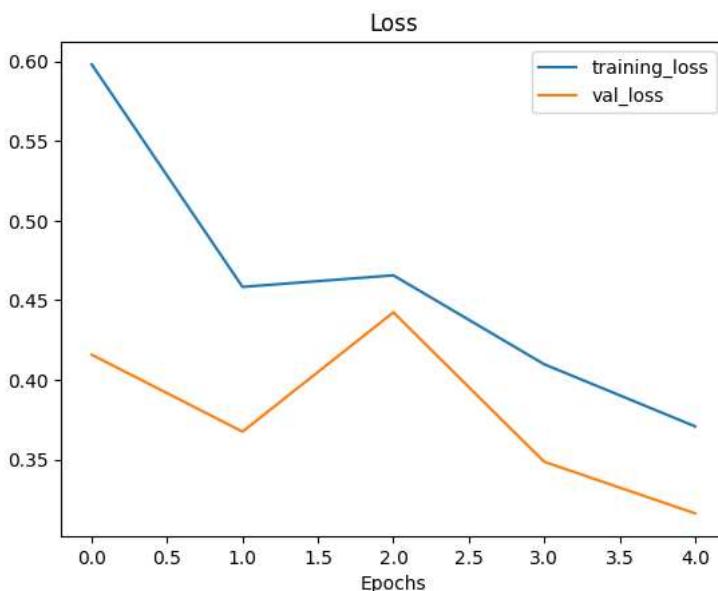
Loss



Accuracy



```
plot_loss_curves(history_1)
```



```
print(class_names)
```

```
['pizza' 'steak']
```

```
!wget https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/images/03-steak.jpeg
steak = mpimg.imread("03-steak.jpeg")
plt.imshow(stake)
plt.axis(False);
```

```
--2023-05-11 04:00:27-- https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/images/03-steak.jpeg
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.108.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1978213 (1.9M) [image/jpeg]
Saving to: '03-steak.jpeg'
```

```
03-steak.jpeg      100%[=====] 1.89M --KB/s in 0.02s
```

```
2023-05-11 04:00:27 (90.1 MB/s) - '03-steak.jpeg' saved [1978213/1978213]
```



```
# Check the shape of our image
steak.shape
```

```
(4032, 3024, 3)
```

```

def load_and_prep_image(filename, img_shape=224):
    """
    Reads an image from filename, turns it into a tensor
    and reshapes it to (img_shape, img_shape, colour_channel).
    """
    # Read in target file (an image)
    img = tf.io.read_file(filename)

    # Decode the read file into a tensor & ensure 3 colour channels
    # (our model is trained on images with 3 colour channels and sometimes images have 4 colour channels)
    img = tf.image.decode_image(img, channels=3)

    # Resize the image (to the same size our model was trained on)
    img = tf.image.resize(img, size = [img_shape, img_shape])

    # Rescale the image (get all values between 0 and 1)
    img = img/255.
    return img

# Load in and preprocess our custom image
steak = load_and_prep_image("03-steak.jpeg")
steak

<tf.Tensor: shape=(224, 224, 3), dtype=float32, numpy=
array([[[0.6377451 , 0.6220588 , 0.57892156],
       [0.6504902 , 0.63186276, 0.5897059 ],
       [0.63186276, 0.60833335, 0.5612745 ],
       ...,
       [0.52156866, 0.05098039, 0.09019608],
       [0.49509802, 0.04215686, 0.07058824],
       [0.52843136, 0.07745098, 0.10490196]],

      [[0.6617647 , 0.6460784 , 0.6107843 ],
       [0.6387255 , 0.6230392 , 0.57598037],
       [0.65588236, 0.63235295, 0.5852941 ],
       ...,
       [0.5352941 , 0.06862745, 0.09215686],
       [0.529902 , 0.05931373, 0.09460784],
       [0.5142157 , 0.05539216, 0.08676471]],

      [[0.6519608 , 0.6362745 , 0.5892157 ],
       [0.6392157 , 0.6137255 , 0.56764704],
       [0.65637255, 0.6269608 , 0.5828431 ],
       ...,
       [0.53137255, 0.06470589, 0.08039216],
       [0.527451 , 0.06862745, 0.1        ],
       [0.52254903, 0.05196078, 0.0872549 ]],

      ...,

      [[0.49313724, 0.42745098, 0.31029412],
       [0.05441177, 0.01911765, 0.        ],
       [0.2127451 , 0.16176471, 0.09509804],
       ...,
       [0.6132353 , 0.59362745, 0.57009804],
       [0.65294117, 0.6333333 , 0.6098039 ],
       [0.64166665, 0.62990195, 0.59460783]],

      [[0.65392154, 0.5715686 , 0.45      ],
       [0.6367647 , 0.54656863, 0.425     ],
       [0.04656863, 0.01372549, 0.        ],
       ...,
       [0.6372549 , 0.61764705, 0.59411764],
       [0.63529414, 0.6215686 , 0.5892157 ],
       [0.6401961 , 0.62058824, 0.59705883]],

      [[0.1        , 0.05539216, 0.        ],
       [0.48333332, 0.40882352, 0.29117647],
       [0.65        , 0.5686275 , 0.44019607],
       ...,
       [0.6308824 , 0.6161765 , 0.5808824 ],
       [0.6519608 , 0.63186276, 0.5901961 ],
       [0.6338235 , 0.6259804 , 0.57892156]]], dtype=float32)>

```

model_8.predict(steak)

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-61-fd7eef5274d1> in <cell line: 2>()
      1 # Make a prediction on our custom image (spoiler: this won't work)
----> 2 model_8.predict(steak)

----- 1 frames -----
/usr/local/lib/python3.10/dist-packages/keras/engine/training.py in
tf__predict_function(iterator)
    13         try:
    14             do_return = True
--> 15         retval_ = ag__.converted_call(ag__.ld(step_function),
(ag__.ld(self), ag__.ld(iterator)), None, fscope)
    16     except:
    17         do_return = False

ValueError: in user code:

  File "/usr/local/lib/python3.10/dist-packages/keras/engine/training.py", line 2169,
in predict_function *
    return step_function(self, iterator)
  File "/usr/local/lib/python3.10/dist-packages/keras/engine/training.py", line 2155,
in step_function **
    outputs = model.distribute_strategy.run(run_step, args=(data,))
  File "/usr/local/lib/python3.10/dist-packages/keras/engine/training.py", line 2143,
in run_step **
    outputs = model.predict_step(data)
  File "/usr/local/lib/python3.10/dist-packages/keras/engine/training.py", line 2111,
in predict_step
    return self(x, training=False)

# Add an extra axis
print(f"Shape before new dimension: {steak.shape}")
steak = tf.expand_dims(steak, axis=0) # add an extra dimension at axis 0
#steak = steak[tf.newaxis, ...] # alternative to the above, '...' is short for 'every other dimension'
print(f"Shape after new dimension: {steak.shape}")
steak

Shape before new dimension: (224, 224, 3)
Shape after new dimension: (1, 224, 224, 3)
<tf.Tensor: shape=(1, 224, 224, 3), dtype=float32, numpy=
array([[[[0.6377451 , 0.6220588 , 0.57892156],
       [0.6504902 , 0.63186276, 0.5897059 ],
       [0.63186276, 0.60833335, 0.5612745 ],
       ...,
       [0.52156866, 0.05098039, 0.09019608],
       [0.49509802, 0.04215686, 0.07058824],
       [0.52843136, 0.07745098, 0.10490196]],

      [[0.6617647 , 0.6460784 , 0.6107843 ],
       [0.6387255 , 0.6230392 , 0.57598037],
       [0.65588236, 0.63235295, 0.5852941 ],
       ...,
       [0.5352941 , 0.06862745, 0.09215686],
       [0.529902 , 0.05931373, 0.09460784],
       [0.5142157 , 0.05539216, 0.08676471]],

      [[0.6519608 , 0.6362745 , 0.5892157 ],
       [0.6392157 , 0.6137255 , 0.56764704],
       [0.65637255, 0.6269608 , 0.5828431 ],
       ...,
       [0.53137255, 0.06470589, 0.08039216],
       [0.527451 , 0.06862745, 0.1        ],
       [0.52254903, 0.05196078, 0.0872549 ]],

      ...,
      [[0.49313724, 0.42745098, 0.31029412],
       [0.05441177, 0.01911765, 0.        ],
       [0.2127451 , 0.16176471, 0.09509804],
       ...,
       [0.6132353 , 0.59362745, 0.57009804],
       [0.65294117, 0.6333333 , 0.6098039 ],
       [0.64166665, 0.62990195, 0.59460783]],

      [[0.65392154, 0.5715686 , 0.45      ],
       [0.6367647 , 0.54656863, 0.425     ],
       [0.04656863, 0.01372549, 0.        ],
       ...,
       [0.6372549 , 0.61764705, 0.59411764],
       [0.63529414, 0.6215686 , 0.5892157 ],
       [0.6401961 , 0.62058824, 0.59705883]]],
```

```
[[0.1      , 0.05539216, 0.        ],
 [0.48333332, 0.40882352, 0.29117647],
 [0.65      , 0.5686275 , 0.44019607],
 ...,
 [0.6308824 , 0.6161765 , 0.5808824 ],
 [0.6519608 , 0.63186276, 0.5901961 ],
 [0.6338235 , 0.6259804 , 0.57892156]]], dtype=float32)>
```

```
# Make a prediction on custom image tensor
pred = model_8.predict(steak)
pred

1/1 [=====] - 0s 401ms/step
array([[0.7315909]], dtype=float32)
```

```
class_names

array(['pizza', 'steak'], dtype='<U5')
```

```
pred_class = class_names[int(tf.round(pred)[0][0])]
pred_class

'steak'
```

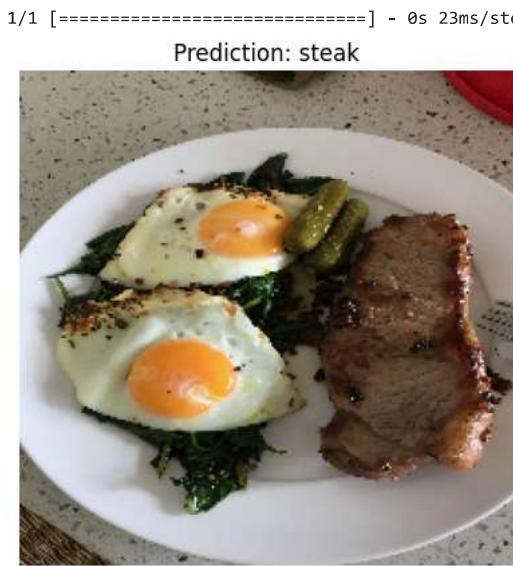
```
def pred_and_plot(model, filename, class_names):
    """
    Imports an image located at filename, makes a prediction on it with
    a trained model and plots the image with the predicted class as the title.
    """
    # Import the target image and preprocess it
    img = load_and_prep_image(filename)

    # Make a prediction
    pred = model.predict(tf.expand_dims(img, axis=0))

    # Get the predicted class
    pred_class = class_names[int(tf.round(pred)[0][0])]

    # Plot the image and predicted class
    plt.imshow(img)
    plt.title(f"Prediction: {pred_class}")
    plt.axis(False);
```

```
pred_and_plot(model_8, "03-steak.jpeg", class_names)
```



▼ 1. Import and become one with the data

Again, we've got a subset of the [Food101 dataset](#). In addition to the pizza and steak images, we've pulled out another eight classes.

```

import os

# Walk through 10_food_classes directory and list number of files
for dirpath, dirnames, filenames in os.walk("10_food_classes_all_data"):
    print(f"There are {len(dirnames)} directories and {len(filenames)} images in '{dirpath}'.")

There are 2 directories and 0 images in '10_food_classes_all_data'.
There are 10 directories and 0 images in '10_food_classes_all_data/test'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/chicken_wings'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/fried_rice'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/ice_cream'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/grilled_salmon'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/chicken_curry'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/hamburger'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/steak'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/pizza'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/sushi'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/ramen'.
There are 10 directories and 0 images in '10_food_classes_all_data/train'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/chicken_wings'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/fried_rice'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/ice_cream'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/grilled_salmon'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/chicken_curry'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/hamburger'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/steak'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/pizza'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/sushi'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/ramen'.

```

```

train_dir = "10_food_classes_all_data/train/"
test_dir = "10_food_classes_all_data/test/"

```

```

# Get the class names
import pathlib
import numpy as np
data_dir = pathlib.Path(train_dir)
class_names = np.array(sorted([item.name for item in data_dir.glob('*')]))
print(class_names)

['chicken_curry' 'chicken_wings' 'fried_rice' 'grilled_salmon' 'hamburger'
 'ice_cream' 'pizza' 'ramen' 'steak' 'sushi']

```

```

# View a random image from the training dataset
import random
img = view_random_image(target_dir=train_dir,
                        target_class=random.choice(class_names)) # get a random class name

```



```

from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Rescale the data and create data generator instances
train_datagen = ImageDataGenerator(rescale=1/255.)
test_datagen = ImageDataGenerator(rescale=1/255.)

```

```
# Load data in from directories and turn it into batches
train_data = train_datagen.flow_from_directory(train_dir,
                                                target_size=(224, 224),
                                                batch_size=32,
                                                class_mode='categorical') # changed to categorical

test_data = train_datagen.flow_from_directory(test_dir,
                                              target_size=(224, 224),
                                              batch_size=32,
                                              class_mode='categorical')

Found 7500 images belonging to 10 classes.
Found 2500 images belonging to 10 classes.

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense

model_9 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(10, activation='softmax') # changed to have 10 neurons (same as number of classes) and 'softmax' activation
])

# Compile the model
model_9.compile(loss="categorical_crossentropy", # changed to categorical_crossentropy
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=["accuracy"])

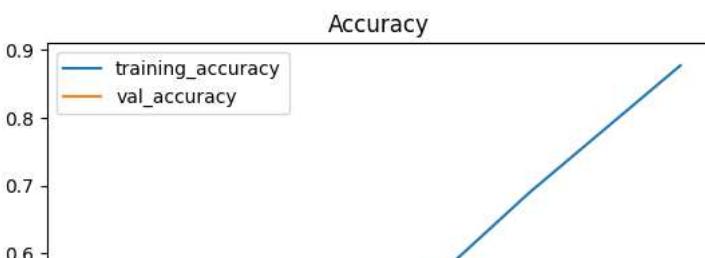
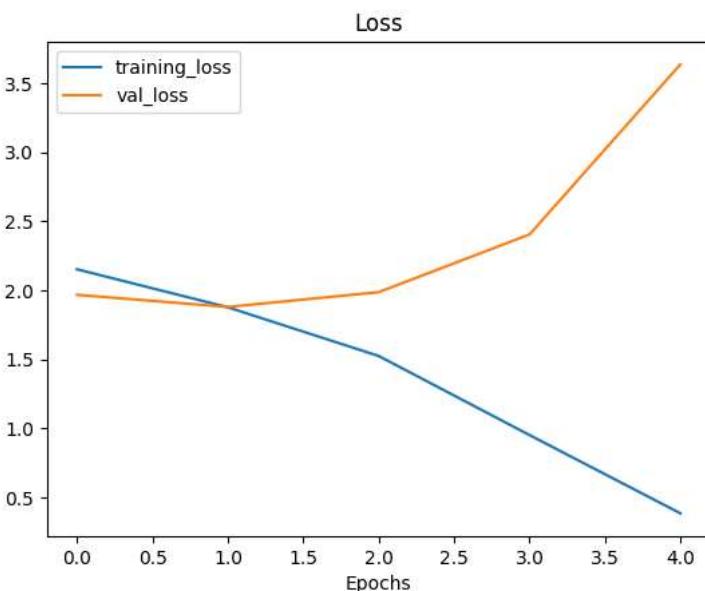
# Fit the model
history_9 = model_9.fit(train_data, # now 10 different classes
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=test_data,
                        validation_steps=len(test_data))

Epoch 1/5
235/235 [=====] - 48s 194ms/step - loss: 2.1520 - accuracy: 0.2133 - val_loss: 1.9674 - val_accuracy: 0.2964
Epoch 2/5
235/235 [=====] - 45s 193ms/step - loss: 1.8770 - accuracy: 0.3489 - val_loss: 1.8792 - val_accuracy: 0.3424
Epoch 3/5
235/235 [=====] - 45s 190ms/step - loss: 1.5239 - accuracy: 0.4860 - val_loss: 1.9865 - val_accuracy: 0.3184
Epoch 4/5
235/235 [=====] - 45s 193ms/step - loss: 0.9507 - accuracy: 0.6897 - val_loss: 2.4049 - val_accuracy: 0.2824
Epoch 5/5
235/235 [=====] - 45s 190ms/step - loss: 0.3835 - accuracy: 0.8772 - val_loss: 3.6373 - val_accuracy: 0.2760

# Evaluate on the test data
model_9.evaluate(test_data)

79/79 [=====] - 11s 139ms/step - loss: 3.6373 - accuracy: 0.2760
[3.6373109817504883, 0.2759999930858612]

# Check out the model's loss curves on the 10 classes of data (note: this function comes from above in the notebook)
plot_loss_curves(history_9)
```



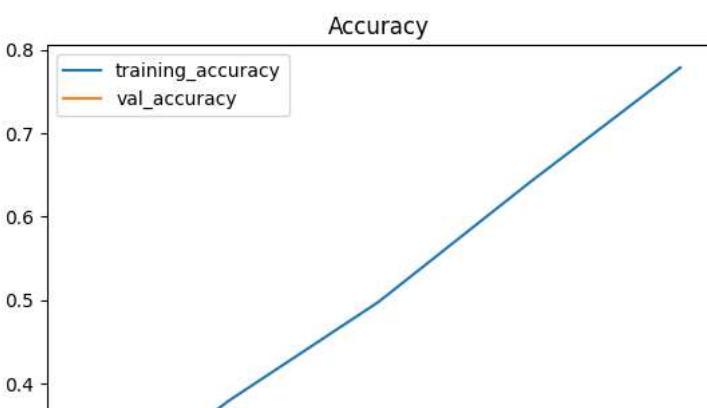
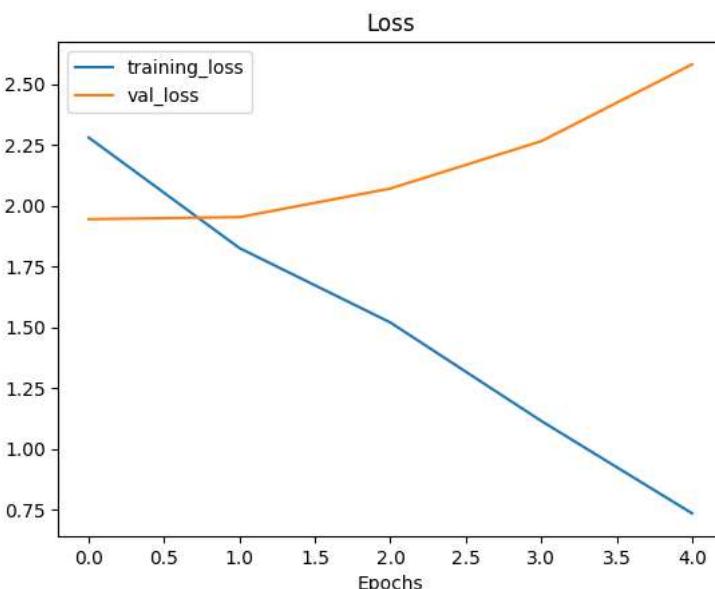
```
# Try a simplified model (removed two layers)
model_10 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(10, activation='softmax')
])

model_10.compile(loss='categorical_crossentropy',
                  optimizer=tf.keras.optimizers.Adam(),
                  metrics=['accuracy'])

history_10 = model_10.fit(train_data,
                           epochs=5,
                           steps_per_epoch=len(train_data),
                           validation_data=test_data,
                           validation_steps=len(test_data))

Epoch 1/5
235/235 [=====] - 46s 193ms/step - loss: 2.2804 - accuracy: 0.2397 - val_loss: 1.9445 - val_accuracy: 0.3024
Epoch 2/5
235/235 [=====] - 45s 192ms/step - loss: 1.8253 - accuracy: 0.3785 - val_loss: 1.9529 - val_accuracy: 0.3172
Epoch 3/5
235/235 [=====] - 44s 186ms/step - loss: 1.5193 - accuracy: 0.4979 - val_loss: 2.0705 - val_accuracy: 0.3160
Epoch 4/5
235/235 [=====] - 44s 187ms/step - loss: 1.1143 - accuracy: 0.6408 - val_loss: 2.2653 - val_accuracy: 0.3064
Epoch 5/5
235/235 [=====] - 44s 188ms/step - loss: 0.7340 - accuracy: 0.7787 - val_loss: 2.5814 - val_accuracy: 0.2944

# Check out the loss curves of model_10
plot_loss_curves(history_10)
```



```
# Create augmented data generator instance
train_datagen_augmented = ImageDataGenerator(rescale=1/255.,
                                              rotation_range=20, # note: this is an int not a float
                                              width_shift_range=0.2,
                                              height_shift_range=0.2,
                                              zoom_range=0.2,
                                              horizontal_flip=True)

train_data_augmented = train_datagen_augmented.flow_from_directory(train_dir,
                                                               target_size=(224, 224),
                                                               batch_size=32,
                                                               class_mode='categorical')
```

Found 7500 images belonging to 10 classes.

```
# Clone the model
model_11 = tf.keras.models.clone_model(model_10)

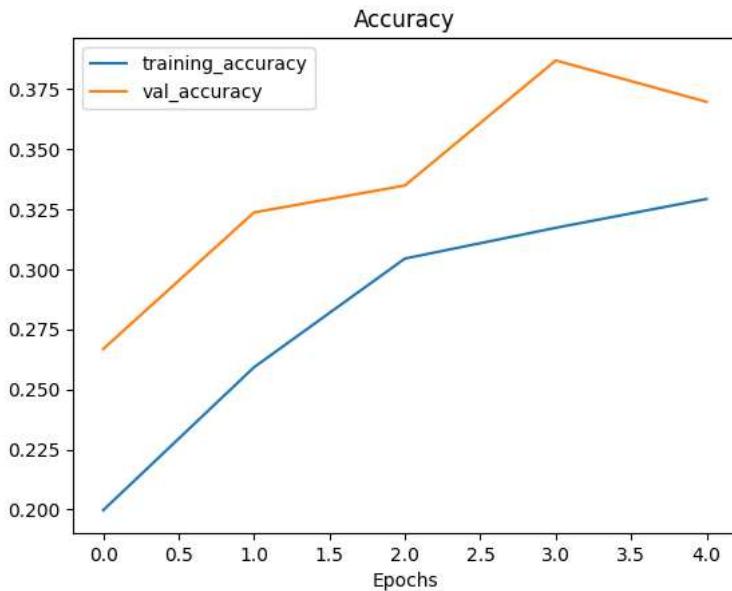
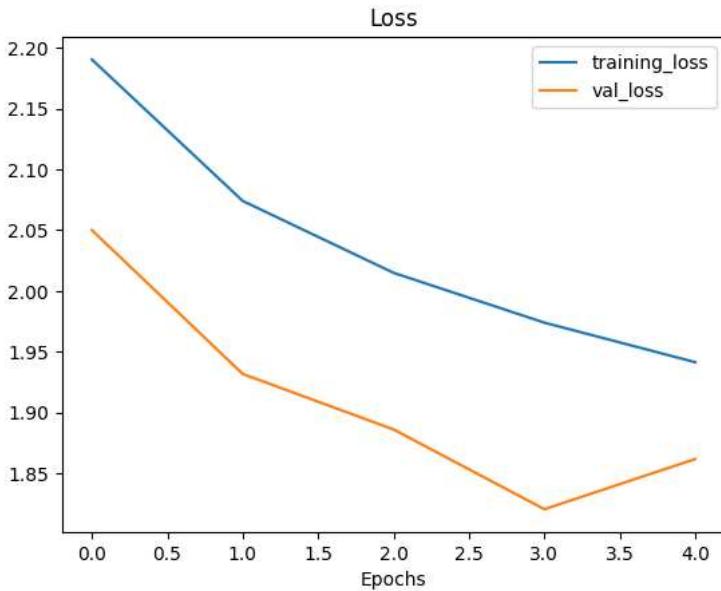
# Compile the cloned model
model_11.compile(loss="categorical_crossentropy",
                  optimizer=tf.keras.optimizers.Adam(),
                  metrics=["accuracy"])

# Fit the model
history_11 = model_11.fit(train_data_augmented, # use augmented data
                           epochs=5,
                           steps_per_epoch=len(train_data_augmented),
                           validation_data=test_data,
                           validation_steps=len(test_data))
```

```
Epoch 1/5
235/235 [=====] - 112s 473ms/step - loss: 2.1903 - accuracy: 0.1997 - val_loss: 2.0500 - val_accuracy: 0.2668
Epoch 2/5
235/235 [=====] - 114s 483ms/step - loss: 2.0740 - accuracy: 0.2592 - val_loss: 1.9318 - val_accuracy: 0.3236
Epoch 3/5
235/235 [=====] - 114s 485ms/step - loss: 2.0149 - accuracy: 0.3044 - val_loss: 1.8862 - val_accuracy: 0.3348
```

```
Epoch 4/5
235/235 [=====] - 114s 487ms/step - loss: 1.9739 - accuracy: 0.3172 - val_loss: 1.8206 - val_accuracy: 0.3868
Epoch 5/5
235/235 [=====] - 115s 489ms/step - loss: 1.9415 - accuracy: 0.3292 - val_loss: 1.8619 - val_accuracy: 0.3696
```

```
# Check model's performance with augmented data
plot_loss_curves(history_11)
```



```
# What classes has our model been trained on?
class_names

array(['chicken_curry', 'chicken_wings', 'fried_rice', 'grilled_salmon',
       'hamburger', 'ice_cream', 'pizza', 'ramen', 'steak', 'sushi'],
      dtype='<U14')

# -q is for "quiet"
!wget -q https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/images/03-pizza-dad.jpeg
!wget -q https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/images/03-steak.jpeg
!wget -q https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/images/03-hamburger.jpeg
!wget -q https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/images/03-sushi.jpeg

# Make a prediction using model_11
pred_and_plot(model=model_11,
              filename="03-steak.jpeg",
              class_names=class_names)
```

```
1/1 [=====] - 0s 81ms/step
```

Prediction: chicken_curry



```
pred_and_plot(model_11, "03-sushi.jpeg", class_names)
```

```
1/1 [=====] - 0s 23ms/step
```

Prediction: chicken_curry



```
# Load in and preprocess our custom image
img = load_and_prep_image("03-steak.jpeg")

# Make a prediction
pred = model_11.predict(tf.expand_dims(img, axis=0))

# Match the prediction class to the highest prediction probability
pred_class = class_names[pred.argmax()]
plt.imshow(img)
plt.title(pred_class)
plt.axis(False);
```

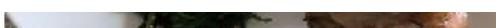
```
1/1 [=====] - 0s 22ms/step
```

steak



```
# Check the output of the predict function
pred = model_11.predict(tf.expand_dims(img, axis=0))
pred
```

```
1/1 [=====] - 0s 22ms/step
array([0.00413127, 0.00698275, 0.00470221, 0.14550138, 0.0163848 ,
       0.00341261, 0.00258247, 0.2162087 , 0.5952408 , 0.00485306]),
      dtype=float32)
```



```
# Find the predicted class name
class_names[pred.argmax()]
```

'steak'

```
# Adjust function to work with multi-class
```

```
def pred_and_plot(model, filename, class_names):
    """
```

```
    Imports an image located at filename, makes a prediction on it with
    a trained model and plots the image with the predicted class as the title.
```

```
    """
    # Import the target image and preprocess it
    img = load_and_prep_image(filename)
```

```
    # Make a prediction
    pred = model.predict(tf.expand_dims(img, axis=0))
```

```
    # Get the predicted class
```

```
    if len(pred[0]) > 1: # check for multi-class
        pred_class = class_names[pred.argmax()] # if more than one output, take the max
    else:
        pred_class = class_names[int(tf.round(pred)[0][0])] # if only one output, round
```

```
    # Plot the image and predicted class
```

```
    plt.imshow(img)
    plt.title(f"Prediction: {pred_class}")
    plt.axis(False);
```

```
pred_and_plot(model_11, "03-steak.jpeg", class_names)
```

```
1/1 [=====] - 0s 23ms/step
```

Prediction: steak



```
pred_and_plot(model_11, "03-sushi.jpeg", class_names)
```

```
1/1 [=====] - 0s 22ms/step
```

Prediction: grilled_salmon



```
pred_and_plot(model_11, "03-hamburger.jpeg", class_names)
```

```
1/1 [=====] - 0s 22ms/step
```

Prediction: steak



```
# Save a model
```

```
model_11.save("saved_trained_model")
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 2 of 2).
```

```
# Load in a model and evaluate it
```

```
loaded_model_11 = tf.keras.models.load_model("saved_trained_model")
loaded_model_11.evaluate(test_data)
```

```
79/79 [=====] - 11s 141ms/step - loss: 1.8619 - accuracy: 0.3696
[1.8618727922439575, 0.36959999799728394]
```

```
model_11.evaluate(test_data)
```

```
79/79 [=====] - 11s 140ms/step - loss: 1.8619 - accuracy: 0.3696
[1.8618724346160889, 0.36959999799728394]
```

